

# GEMM3: Constant-Workspace High-Performance Multiplication of Three Matrices for Matrix Chaining

Krzysztof A. Drewniak

The University of Texas at Austin

April 13, 2018

# Matrix chaining problem

- ▶ Problem: compute  $A_1 A_2 \cdots A_N$  efficiently
- ▶  $O(N \log N)$  algorithm<sup>1</sup>, also  $O(N^3)$  with dynamic programming<sup>2</sup>
- ▶ Fewer flops  $\rightarrow$  more performance?
- ▶ Ex:  $WXYZ \Rightarrow ((WX)Y)Z, (WX)(YZ), W(X(YZ)), \dots$

---

<sup>1</sup>Hu and Shing, 1984

<sup>2</sup>Barthels 2018

# Generalized matrix chaining

- ▶ In reality — transposes, inverses, properties
- ▶ Ensemble Kalman filter<sup>3</sup>  $X_i^b S_i (Y_i^b)^T R_i^{-1}$   
 Tridiagonalization<sup>4</sup>  $\tau_u \tau_v v v^T A u u^T$   
 Two-sided triangular solve<sup>5</sup>  $L^{-1} A L^{-H}$  ( $L$  lower triangular)
- ▶ Performance with BLAS/LAPACK<sup>6</sup> – must be expert
- ▶ Less performance with Matlab, numpy, etc. (left-to-right)
- ▶ Linnea<sup>7</sup>: expression  $\rightarrow$  BLAS calls automatically

---

<sup>3</sup>Rao 2017

<sup>4</sup>Choi 1995

<sup>5</sup>Poulson 2011

<sup>6</sup>Dongarra 1990, Anderson 1999

<sup>7</sup>Barthels 2018

# GEMM3 — Why bother?

- ▶ ▶  $\mathbf{X}_i^b \mathbf{S}_i (\mathbf{Y}_i^b)^T \mathbf{R}_i^{-1}$
- ▶ ▶  $\tau_u \tau_v \mathbf{v} \mathbf{v}^T \mathbf{A} \mathbf{u} \mathbf{u}^T$
- ▶ ▶  $\mathbf{L}^{-1} \mathbf{A} (\mathbf{L}^{-1})^H$  ( $\mathbf{L}$  lower triangular)
- ▶ All multiply three matrices as a subproblem
- ▶ Not all problems subdivide like this

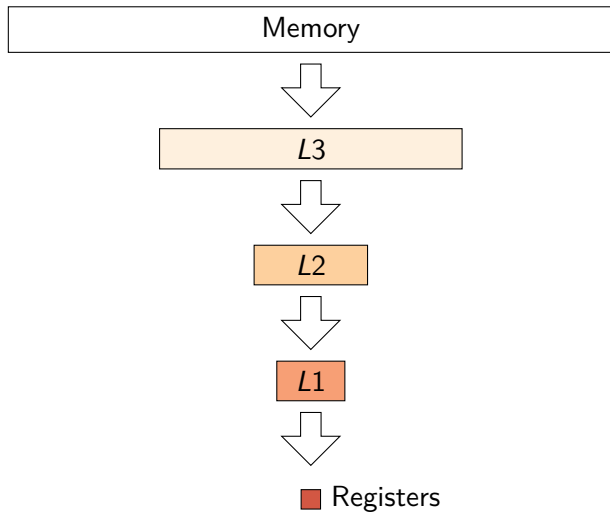
## GEMM3 — Why a new algorithm?

- ▶ Current approach: parentheses, multiply twice, store temporary  $T$
- ▶  $T$  often eats memory (& performance)
- ▶ We can do better!
- ▶ Use how GEMM works to nest computations
- ▶  $O(1)$  extra memory, maybe more performance

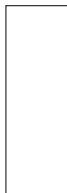
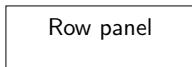
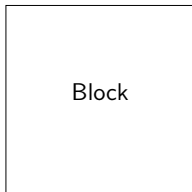
## Section 2

# High-Performance GEMM

# Memory hierarchy



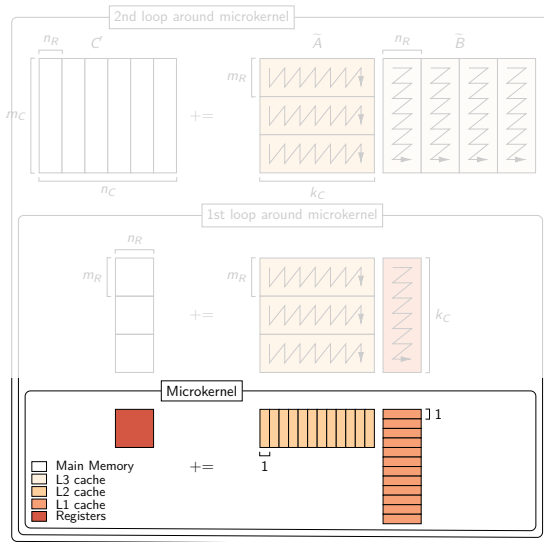
# Important matrix shapes



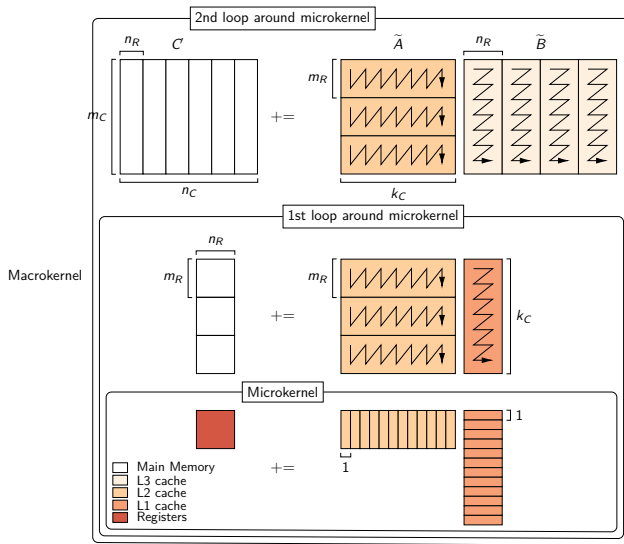
Column panel



# GEMM: The kernels



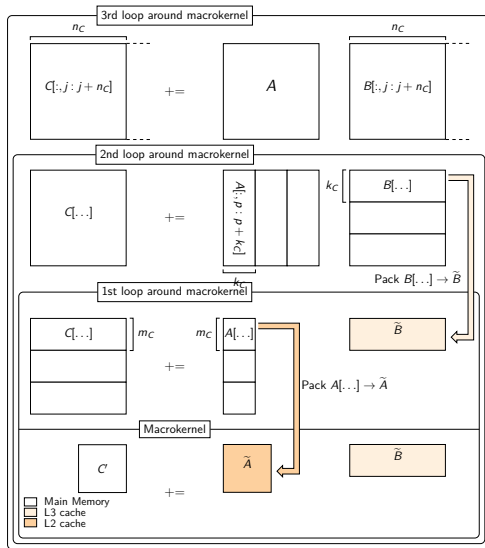
# GEMM: The kernels



Packing very important<sup>8</sup>

<sup>8</sup>Henry 1992

# GEMM: The algorithm<sup>9</sup>

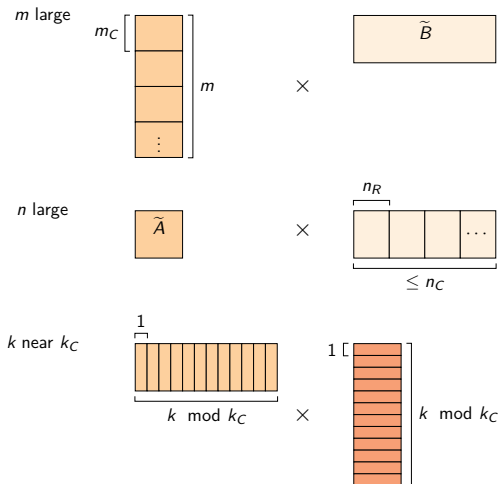


<sup>9</sup>Goto 2008

# Data reuse

- Every loop reads *something* repeatedly

Want:

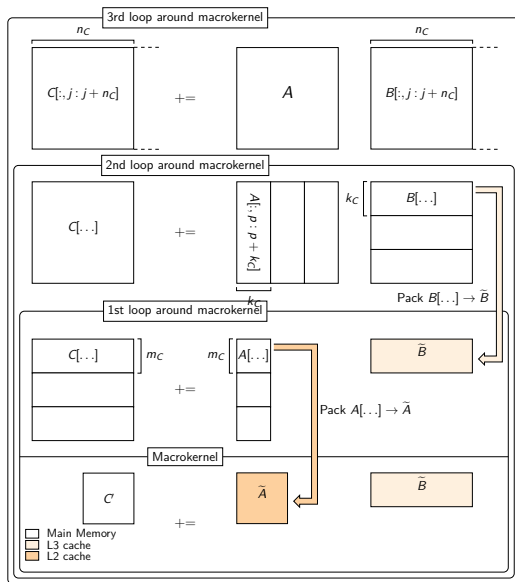


## Section 3

# The GEMM3 algorithm

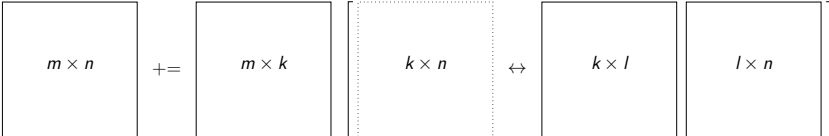
# Key concept of the algorithm

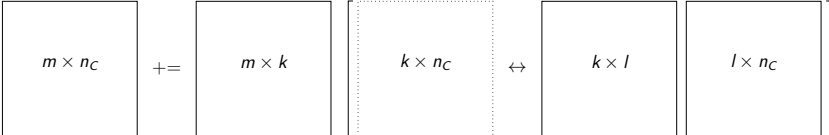
- ▶ We want  $G += D(EF)$ , (dimensions:  $m, k, l, n$  in order)
- ▶  $EF$  first needed in packing step
- ▶ Don't do computation until then

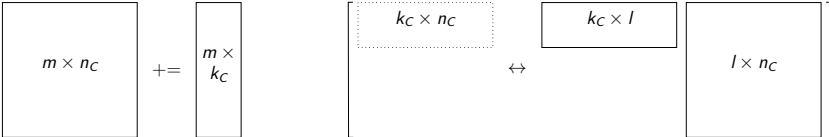


# Deriving GEMM3: Partitionings

$$G \quad += \quad D \quad [(EF) \quad \leftrightarrow \quad E \quad F]$$

1. 

2. 

3. 

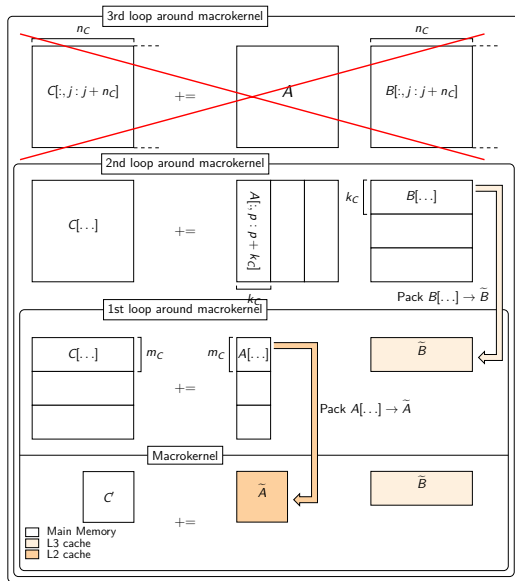


# Deriving GEMM3: Inner algorithm


$$\boxed{EF: k_C \times n_C} = \boxed{E: k_C \times l} \boxed{F: l \times n_C}$$

- ▶ Only point to compute  $EF$  in constant memory
- ▶ GEMM algorithm needs tweaks


# Inner algorithm tweaks: Removing the outer loop




# Inner algorithm tweaks: Microkernel packed writes




# Inner algorithm tweaks: Halving $n_C$




# Inner algorithm tweaks: Small $k_C$ reduction



# Inner algorithm tweaks: Small $k_C$ reduction



# The algorithm



# The small drawback

Problem shape:

$$\widetilde{EF}: k_C \times n_C = E: k_C \times l \quad F: l \times n_C$$

Reuse problem:  $m$  small

$$m_C \left[ \begin{array}{c} \text{orange box} \\ \text{orange box} \\ \text{crossed out orange box} \\ \text{crossed out orange box} \end{array} \right] k_C \times \widetilde{F}$$



## Section 4

# Experiments and Results

## Implementation details

- ▶ Multilevel Optimization of Matrix Multiply Sandbox (MOMMS)<sup>10</sup>
- ▶ Extended to support three matrices
- ▶ Implement both GEMM3 and pair of GEMM algorithms
- ▶ GEMM (from BLIS<sup>11</sup>) port performs like BLIS
- ▶ Machine: 3.5 GHz (one core used), 15 GB RAM, 32 KB  $L1$  cache, 256 KB  $L2$ , 8 MB  $L3$ . Peak perf 56 GFLOPS/s.

	GEMM3	GEMM algorithm
$m_R$	6	6
$n_R$	8	8
$m_C$	72	72
$k_C$	252	256
$l_C$	256	
$n_C$	2040	4080

Table: Parameters for Haswell CPUs

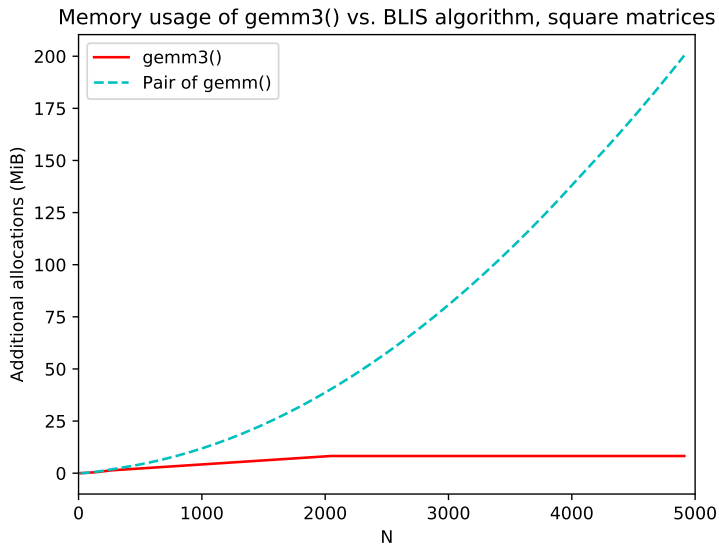
<sup>10</sup>Smith 2018

<sup>11</sup>van Zee 2016

# Experiments

1.  $G += D(EF)$ , square matrices
  - ▶ Inputs column-major, outputs row-major for fairness
2.  $G^T += F^T(E^T D^T)$ , square matrices
  - ▶ After transpose, all row major
3.  $G += D(EF)$ , rectangles (one dimension small)

# Workspace usage, square matrices




## API simplicity

```
double *T = malloc(k * n * sizeof(double));
dgemm("N", "N", k, l, n,
      1, E, lde, F, ldf,
      0, T, k);
dgemm("N", "N", m, k, n,
      alpha, D, ldd, T, k,
      beta, G, ldg);
free(T);
```


vs.

```
dgemm3("R", "N", "N", "N", m, k, l, n,
      alpha, D, ldd, E, lde, F, ldf,
      beta, G, ldg);
```

# $G += D(EF)$ , square matrices




# $G += D(EF)$ , square matrices



$$G += (DE)F$$

Ex: **Don't:**




- ▶ Putting parentheses there sometimes better
- ▶ Deriving directly doesn't work — bad shape
- ▶ However,  $G += (DE)F \Leftrightarrow G^T += F^T(E^T D^T)$




$$G += (DE)F$$

Ex: **Do:**

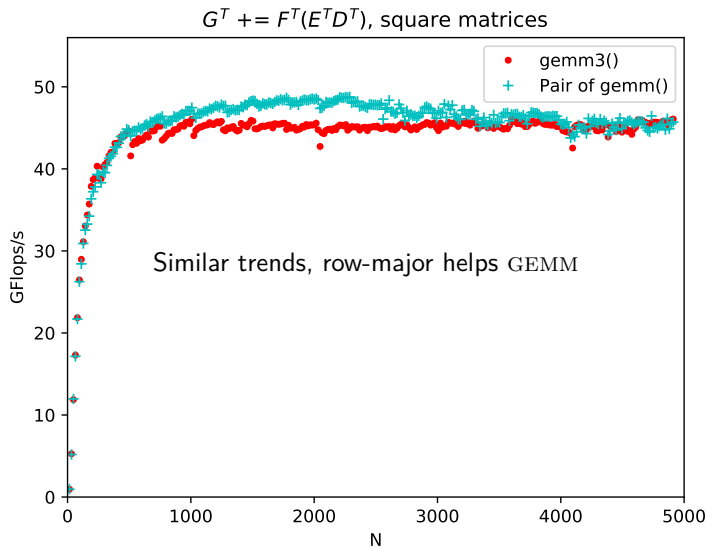


- ▶ Putting parentheses there sometimes better
- ▶ Deriving directly doesn't work — bad shape
- ▶ However,  $G += (DE)F \Leftrightarrow G^T += F^T(E^T D^T)$

# $G += (DE)F$ , square matrices




# $G += (DE)F$ , square matrices



# $G += D(EF)$ , rectangular matrices

$G += D(EF)$ , narrow dimension = 9



# Conclusions

- ▶ GEMM structure lets us make GEMM3
- ▶ Constant memory
- ▶ Cleaner API
- ▶ Comparable performance

# Future Work

- ▶ Parallel case
- ▶ More architectures
- ▶ Variants (matrices with properties), autogeneration

# Acknowledgments

- ▶ Prof. Robert van de Geijn — advising and providing inspiration
- ▶ Dr. Tyler Smith — writing MOMMS and algorithm design
- ▶ Prof. Tze Meng Low — performance fixes
- ▶ NSF awards CCF-1714091 and ACI-1550493 — funding

# Questions?



## Picking parameters: $m_R, n_R$

- ▶ Determine microkernel
- ▶ Based on microarchitecture — register width, FMA properties
- ▶ We're reusing BLIS's work
- ▶ Can swap  $m_R$  and  $n_R$

## Picking parameters: $k_C$

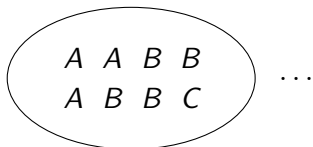
Placing memory in cache: [tag][set #][offset in line]

$$m_r k_C S_{elem} = C_A C_{L1} N_{L1}$$

$$n_r k_C S_{elem} = C_B C_{L1} N_{L1}$$

L1

Cache:



$$C_A + C_B + 1 \leq W_{L1}$$

Maximizing  $k_C$  improves performance

$$C_B = \left\lceil \frac{n_R k_C S_{elem}}{N_{L1} C_{L1}} \right\rceil$$

$$= \left\lceil \frac{n_R}{m_R} C_A \right\rceil$$

$$C_A \leq \left\lfloor \frac{W_{L1} - 1}{1 + \frac{n_R}{m_R}} \right\rfloor$$

GEMM3

## Picking parameters: $m_C$ and $n_C$

- ▶ For  $m_C$ : reserve ways for  $B$  and  $C$
- ▶ Then take all you can
- ▶  $n_C$ , leave out what architecture requires, then divide
- ▶ L3 is very big, tuning is much less needed